

- Homework 4 Spec sheet -

In this homework, we are revisiting a dataset you are already familiar with, with new methods. The learning goal is to realize what these methods can yield, relative to more classic methods. So please load and use the “diabetes.csv” data file.

This dataset contains information from over 250,000 people obtained by the CDC.

The first row represents the column header.

Each row after that represents the information of one person.

Columns represent (in order):

- 1) Diabetes status (1 = has been diagnosed with diabetes, 0 = has not)
- 2) High blood pressure (1 = has been diagnosed with hypertension, 0 = has not)
- 3) High cholesterol (1 = has been diagnosed with high cholesterol, 0 = has not)
- 4) Body Mass Index (weight / height²)
- 5) Smoker (1 = person has smoked more than 100 cigarettes in their life, 0 = has not)
- 6) Stroke (1 = person has previously suffered a stroke, 0 = has not)
- 7) Myocardial issues (1 = has previously had a heart attack, 0 = has not)
- 8) Physically active (1 = person describes themselves as physically active, 0 = does not)
- 9) Eats fruit (1 = person reports eating fruit at least once a day, 0 = does not)
- 10) Eats vegetables (1 = person reports eating vegetables at least once a day, 0 = does not)
- 11) Heavy Drinker (1 = consumes more drinks than the CDC threshold/week, 0 = does not)
- 12) Has healthcare (1 = person has some kind of healthcare plan coverage, 0 = does not)
- 13) NotAbleToAffordDoctor (1 = person needed to see the doctor within the last year, but could not afford to, 0 = did not)
- 14) General health: Self-assessment of health status on a scale from 1 to 5
- 15) Mental health: Days of poor mental health in the last 30 days (self-assessed)
- 16) Physical health: Days of poor physical health in the last 30 days (self-assessed)
- 17) Hard to climb stairs (1 = person reports difficulties in climbing stairs, 0 = does not)
- 18) Biological sex (1 = male, 2 = female)
- 19) Age bracket (1 = 18-24, 2 = 25-29, 3 = 30-34, 4 = 35-39, 5 = 40-44, 6 = 45-49, 7 = 50-54, 8 = 55-59, 9 = 60-64, 10 = 65-69, 11 = 70-74, 12 = 75-79, 13 = 80+)
- 20) Education bracket (terminal education is 1 = only kindergarten, 2 = elementary school, 3 = some high school, 4 = GED, 5 = some college, 6 = college graduate)
- 21) Income bracket (Annual income where 1 = below \$10k, 8 = above \$75k)
- 22) Zodiac sign (Tropical calendar, 1 = Aries, 12 = Pisces, with everything else in between)

As before, this data is carefully curated, so there should not be too much missing data (if any). Diabetes is major metabolic disorder, with over 400 million people suffering from this illness worldwide. It is also a leading cause of death, with over 1.5 million annual deaths. It has been suggested that a large proportion of these are due to lifestyle choices. Therefore, for both prevention and treatment, it is critically important to find good predictors of diabetes. We will attempt to do so in this exercise.

Mission command approach: As per §4.5 of the Sittyba, we will tell you what to do (“answer these questions”), not how to do it. That is up to you. However, we want you to:

- a) Do the homework yourself. Do not copy answers from someone else.
- b) Restrict your methods (for now) to what was covered in the lecture/lab (in other words, Perceptron, Feedforward neural network, CNN))
- c) Include the following elements in your answer (so we can grade consistently):

Each answer should contain these elements:

- 1) A brief statement (~paragraph) of what was done to answer the question (narratively explaining what you did in code to answer the question, at a high level).
- 2) A brief statement (~paragraph) as to why this was done (why the question was answered in this way, not by doing something else. Some kind of rationale as to why you did x and not y or z to answer the question – why is what you did a suitable approach?). As there are directed questions here (e.g. “build a feedforward network”, you might want to justify your specific design choices, e.g. number of neurons per layer, cost function, activation functions, etc.)
- 3) A brief statement (~paragraph) as to what was found. This should be as objective and specific as possible – just the results/facts. Do make sure to include numbers and a figure (=a graph or plot) in your statement, to substantiate and illustrate it, respectively.
- 4) A brief statement (~paragraph) as to what you think the findings mean. This is your interpretation of your findings and should answer the original question.

Note: Brief actually means “brief”. There is no need to write a novel. There is value to being concise. A couple of pages should be sufficient for the entire report. Do – however – write a report. A data and code-dump is not very useful or valuable in practice. People who pay you so they can ask you questions usually want them answered. Succinctly. In other words, people like their answers clear, concise and coherent. That’s where the added value is. That’s what is valuable. Less is not worthless, but definitely worth less.

Hint: The dataset is small enough that the specific random seed and the specific train/test split might matter for your specific answer. If that is the case, try different seeds and different splits and make sure to comment on that in your report.

Please answer the following questions in your report:

1. Build and train a Perceptron (one input layer, one output layer, no hidden layers and no activation functions) to classify diabetes from the rest of the dataset. What is the AUC of this model?
2. Build and train a feedforward neural network with at least one hidden layer to classify diabetes from the rest of the dataset. Make sure to try different numbers of hidden layers and different activation functions (at a minimum reLU and sigmoid). Doing so: **How does AUC vary as a function of the number of hidden layers and is it dependent on the kind of activation function used (make sure to include “no activation function” in your comparison). How does this network perform relative to the Perceptron?**
3. Build and train a “deep” network (at least 2 hidden layers) to classify diabetes from the rest of the dataset. Given the nature of this dataset, is there a benefit of using a CNN for the classification?
4. Build and train a feedforward neural network with one hidden layer to predict BMI from the rest of the dataset. Use RMSE to assess the accuracy of your model. Does the RMSE depend on the activation function used?
5. Build and train a neural network of your choice to predict BMI from the rest of your dataset. How low can you get RMSE and what design choices does RMSE seem to depend on?

Extra credit:

- a) Are there any predictors/features that have effectively no impact on the accuracy of these models? If so, please list them and comment briefly on your findings
- b) Write a summary statement on the overall pros and cons of using neural networks to learn from the same dataset as in the prior homework, relative to using classical methods (logistic regression, SVM, trees, forests, boosting methods). Any overall lessons?